Dear Halden Pay Hiring Team,

I am applying for the Senior Data Engineer – Data Platform role. I have nine years of experience building and operating data platforms across e-commerce and payments, with a focus on reusable pipelines, reliable data quality, governed data contracts and production cloud data warehouses that teams can trust. Halden Pay's emphasis on shared platform capabilities and taking ownership of work that helps clients move efficiently strongly appeals to me.

In my current role at Lukovac Freight Hub, I design, build and operate a production data platform at scale, owning its ingestion layer and approximately 140 Airflow DAGs feeding Snowflake. I migrated approximately 300 stored procedures into dbt models, with more than 420 tested models, reducing the nightly warehouse build from 5 hours 10 minutes to 1 hour 35 minutes. I have also introduced data contracts with CI checks, built Kafka-based streaming for around 30 million events a day, and replaced full-table extracts with Debezium CDC. The data-contract work established enforceable data-quality standards between product and data teams, reducing broken-pipeline incidents from approximately nine per month to two. These projects have involved operating production systems, improving performance and reliability, and making platform changes that benefit teams beyond a single pipeline.

My experience also includes payments data: at Kovalent Pay, I built PySpark jobs on AWS EMR to reconcile roughly 1.2 TB of daily raw transaction files, designed a Redshift cloud data warehouse using dimensional modelling, with more than 30 fact and dimension tables, and added data quality checks for critical finance tables. Across my roles, I have applied dimensional modelling, incremental processing and query and warehouse optimisation to large datasets, including a 1.2 TB-per-day reconciliation workload and a Snowflake warehouse whose nightly build time I reduced from 5 hours 10 minutes to 1 hour 35 minutes. The reconciliation engine reduced the finance team's manual work from around three days to four hours per month-end. I bring hands-on experience with Python, SQL, Airflow, dbt, Docker, CI/CD and AWS, as well as Snowflake, Redshift and PySpark, alongside open-source contributions to an Airflow provider, including a fix and tests merged in two pull requests. I would welcome the opportunity to apply this experience with comparable cloud platforms and contribute to Halden Pay's shared data platform through reusable frameworks, data-quality standards and open-source practices.

I would be glad to discuss how I could help Halden Pay build a reliable, reusable platform for its data teams and business. Thank you for your consideration.

Kind regards,
Lena Varga-Holm