Dear Hiring Manager,

I am applying for the Senior Data Engineer – Data Platform role at Halden Pay. I have nine years of experience designing, building and operating production data platforms and cloud warehouses, with hands-on experience across Python, SQL, Airflow, dbt, AWS, Snowflake and Redshift. I am particularly interested in the role's focus on a shared platform and reusable frameworks for analytics, data science and product teams.

In my current role, I own the ingestion layer of a platform with approximately 140 Airflow DAGs and migrated more than 420 transformation models to dbt, adding tests to every model and reducing nightly runtime from 5 hours 10 minutes to 1 hour 35 minutes. I also introduced JSON Schema data contracts enforced through CI, helping reduce broken-pipeline incidents from approximately nine to two per month. These experiences align closely with Halden Pay's focus on dependable platform operations, data quality and data contracts.

My previous work includes building PySpark jobs on AWS EMR to process approximately 1.2 TB of daily transaction files, designing a Redshift warehouse using dimensional/star-schema modelling with more than 30 fact and dimension tables, and developing a reusable internal Python library for S3 and Glue Catalog access, adopted by four teams. I have also built a Kafka-based pipeline handling approximately 30 million shipment-tracking events per day. Alongside this hands-on work, I mentor two engineers and lead weekly platform office hours—experience I would bring to a collaborative team working across disciplines.

I would welcome the opportunity to discuss how my background in scalable data platforms, data modeling and reusable engineering tools could contribute to Halden Pay's Data Platform team. Thank you for your consideration.

Sincerely,

Lena Varga-Holm