Dear Halden Pay Data Platform Team,

I am applying for the Senior Data Engineer – Data Platform role. As a data scientist with over six years in analytics and data science, I bring experience in production SQL, dbt modelling and scheduled forecasting workflows. I am keen to develop further on the production side of data and machine learning, and Halden Pay's focus on shared platforms and reusable tools appeals to me.

At Thistlebox, I write SQL in BigQuery and maintain dbt models for the analytics layer. I built a weekly demand forecasting workflow in Python and LightGBM, using a scheduled Vertex AI Workbench notebook to write forecasts to BigQuery for operations to view in Looker. The model reduced forecast error from 23% to 14% and estimated food waste by 9%. I also introduced pytest unit tests for the forecasting feature code, reaching approximately 60% coverage.

Working with product and growth teams, I built a Streamlit sequential testing calculator now used by around 25 colleagues. I wrote its Dockerfile, while our data platform team deployed and hosts the app on Cloud Run. Presenting results to leadership and delivering internal statistics training have helped me share technical work with colleagues across teams.

I would welcome the opportunity to discuss how my experience with BigQuery, Python and dbt could contribute to Halden Pay's Data Platform team. Thank you for considering my application.

Kind regards,
Tobias Mwangi-Hartley