Dear Halden Pay Data Platform Hiring Team,

I am applying for the Senior Data Engineer – Data Platform role. Over six years in analytics and data science, I have built practical, dependable data workflows using Python, SQL, BigQuery and Google Cloud tooling. In my current role at Thistlebox, I work with Python, production SQL in BigQuery and dbt models, supporting teams that rely on data for operational and product decisions.

I developed a scheduled six-week-ahead demand forecasting workflow in a Vertex AI Workbench notebook using LightGBM, writing forecasts to BigQuery for operations reporting in Looker. It reduced WAPE from 23% to 14% and was associated with an estimated 9% reduction in food waste. I also developed a churn model scored weekly in BigQuery, and built a Streamlit application for planning and analysing sequential experiments. I built a Streamlit application for planning and analysing sequential experiments, wrote its Dockerfile, and worked with the data platform team, which handles its deployment to Cloud Run.

I value reliable, maintainable work: since 2025, I have added pytest unit tests for forecasting feature code, reaching approximately 60% coverage, and I maintain dbt models for the analytics layer and write production SQL in BigQuery. My experience working with product and growth teams, operations, leadership, and a data platform team has taught me to understand the needs behind data requests and build useful solutions. Halden Pay's focus on shared platform frameworks, collaboration and taking ownership strongly appeals to me, and I would welcome the opportunity to contribute my Python, SQL, BigQuery, dbt and data-workflow experience to the Data Platform team.

Thank you for considering my application. I would welcome the opportunity to discuss how my experience could contribute to Halden Pay.

Kind regards,
Tobias Mwangi-Hartley