Dear Hiring Team at Halden Pay,

I am applying for the Senior Data Engineer – Data Platform role. I bring hands-on experience owning data-platform ingestion and orchestration, operating high-volume production pipelines, and improving warehouse reliability and efficiency.

At Lukovac Freight Hub, I own ingestion and orchestration across Airflow, Snowflake, dbt, and Kafka. I built a shipment event pipeline processing around 30 million events per day, with p95 end-to-end latency under two minutes, and added dead-letter and replay tooling to support reliable recovery. I also replaced hourly full-table extracts with Debezium CDC, reducing replica CPU during loads from approximately 70% to 25%. These projects have involved the day-to-day operational work of a production data platform, including an on-call rotation.

Your emphasis on reusable platform capabilities rather than one-off pipelines matches my experience migrating approximately 300 stored procedures into tested and documented dbt models, adding incremental models and slim CI, and reducing the nightly warehouse build from 5 hours 10 minutes to 1 hour 35 minutes. At Kovalent Pay, I built PySpark reconciliation jobs that reduced the finance team's manual month-end work from approximately three days to four hours. That work also connects directly to Halden Pay's cross-border business payments focus: reliable, timely data is essential to reconciliation.

I also contribute to the data engineering community through two merged fixes to an Airflow provider, SQL teaching at a local women-in-tech meetup, and a 2024 talk on data contracts. I would welcome the opportunity to discuss how my platform, data quality, and operational experience could contribute to Halden Pay's Data Platform team.

Kind regards,

Lena Varga-Holm