LENA VARGA-HOLM

Senior Data Engineer

lena.vargaholm@example.com | +385 55 555 0101 | github.com/lvh-data-example

PROFESSIONAL SUMMARY

Data engineer with 9 years of experience building and operating batch and streaming data platforms across logistics, payments and retail. Owns production ingestion and orchestration, with a focus on reusable tooling, tested transformations and reliable data contracts. Combines Python, SQL, Airflow and dbt expertise with AWS cloud experience, dimensional modelling and hands-on platform operations.

TECHNICAL SKILLS

Languages: Python, SQL, Bash
Orchestration and transformation: Airflow, dbt, Spark (PySpark), SSIS
Warehouses and databases: Snowflake, Redshift, PostgreSQL, MySQL, SQL Server
Streaming and CDC: Kafka, Kafka Connect, Debezium, Snowpipe Streaming, AWS MSK
Cloud and engineering: AWS (S3, EMR, Glue, Lambda, IAM), Docker, Git, GitHub Actions, CI/CD, Terraform (basic)
Data quality and modelling: Great Expectations, JSON Schema, data contracts, dimensional modelling, incremental models, pytest
BI: Power BI, Looker (basic)

PROFESSIONAL EXPERIENCE

Senior Data Engineer | Lukovac Freight Hub d.o.o.
03/2022 - present

- Own ingestion and orchestration for the data platform: approximately 140 Airflow DAGs loading PostgreSQL, MySQL, Salesforce and third-party API data into Snowflake; participate in an on-call rotation one week in six.
- Migrated approximately 300 stored procedures into dbt; the platform now has 420+ models with tests on every model. Introduced incremental models, documentation and slim CI in GitHub Actions, cutting nightly build time from 5h10 to 1h35.
- Introduced data contracts using JSON Schema and CI checks between product teams and the data team, reducing broken-pipeline incidents from approximately 9 to 2 per month.
- Designed and built a Kafka and Snowpipe Streaming pipeline processing approximately 30 million shipment events daily at under 2 minutes p95 end-to-end latency; added dead-letter handling and replay tooling.
- Replaced hourly full-table PostgreSQL extracts with Debezium CDC into Kafka, reducing replica CPU during loads from approximately 70% to 25%.
- Reduced monthly Snowflake costs from approximately 38,000 to 24,000 through warehouse sizing, auto-suspend and removal of unused tables in 2024.
- Mentor two mid-level data engineers and run weekly data platform office hours.

Data Engineer | Kovalent Pay
09/2019 - 02/2022

- Built PySpark transaction reconciliation jobs on AWS EMR processing approximately 1.2 TB of raw S3 files daily; reduced finance's manual month-end reconciliation from approximately 3 days to 4 hours.
- Designed a dimensional Redshift warehouse with 30+ fact and dimension tables for finance and risk; implemented Great Expectations checks on critical finance tables.
- Developed a reusable Python library for S3 and Glue Data Catalog access adopted by four teams.

Junior Data Engineer / BI Developer | Adria Retail Group
06/2017 - 08/2019

- Maintained SSIS packages and T-SQL ETL for an on-premises SQL Server warehouse; migrated 60 nightly SSIS jobs to Python scripts scheduled with cron.
- Built Power BI reports for store operations across 40+ stores.

EDUCATION

MSc Computer Science (Data Science track) | Northgate Technical University | 2017
Thesis: Incremental materialized views for analytical queries

CERTIFICATION

AWS Certified Data Engineer - Associate | May 2024

OPEN SOURCE AND COMMUNITY

- Contributed two merged pull requests to an Airflow provider package, fixing a pagination bug and adding tests with Python and pytest (2023).
- Presented "Data contracts without the hype" at a regional data meetup (2024).
- Teach beginner SQL at a local women-in-tech meetup, approximately four evenings per year since 2021.

LANGUAGES

Croatian (native), English (C1), German (B1)