Dear Hiring Team,

I am applying for the Senior Data Engineer – Data Platform role at Halden Pay. My experience combines production data-platform ownership with payment-data engineering, including building reliable pipelines and improving the operational performance of the systems teams depend on.

At Lukovac Freight Hub, I own ingestion and orchestration across Airflow, Snowflake, dbt and Kafka. I designed a shipment-event pipeline processing around 30 million events per day at under two minutes' end-to-end latency at p95, and added dead-letter and replay tooling to support recovery without data loss. I also migrated approximately 300 stored procedures to tested and documented dbt models, cutting the nightly warehouse build from 5 hours 10 minutes to 1 hour 35 minutes. These projects reflect the reusable, well-operated platform approach described in your role.

Earlier, at Kovalent Pay, I built PySpark reconciliation jobs that matched bank settlement files against internal ledgers and reduced manual month-end work for finance from around three days to around four hours. I have also reduced Snowflake costs through warehouse optimisation, contributed fixes and tests to an open-source Airflow provider, and currently mentor two engineers. I would bring this combination of platform ownership, data quality, operational care and payment-domain experience to Halden Pay's Data Platform team.

I am particularly interested in the team's focus on shared platform capabilities and reusable frameworks. I would welcome the opportunity to discuss how my experience could support the data platform and the teams that rely on it.

Kind regards,

Lena Varga-Holm