Dear Hiring Team,

I am applying for the Senior Data Engineer – Data Platform role at Halden Pay. The opportunity to build and operate a shared platform with reusable foundations for analytics, data science, and product teams fits closely with the work I have been doing across data ingestion, orchestration, and warehouse transformation.

At Lukovac Freight Hub, I own platform ingestion and orchestration across Airflow, Snowflake, dbt, and Kafka. I built a shipment-event pipeline processing approximately 30 million events per day, with p95 end-to-end latency under two minutes, and added dead-letter and replay tooling to support recovery without data loss. I also replaced hourly full-table extracts with Debezium CDC, reducing replica CPU during loads from approximately 70% to 25%. Alongside these platform projects, I migrated approximately 300 stored procedures to tested and documented dbt models, cutting the nightly warehouse build from 5 hours 10 minutes to 1 hour 35 minutes.

My earlier work at Kovalent Pay included a PySpark reconciliation engine matching bank settlement files with internal ledgers, reducing month-end manual reconciliation from approximately three days to four hours. Across my roles, I have paired delivery with operational ownership, data quality practices, cost optimisation, mentoring, and open-source contributions. My experience is primarily with AWS and Snowflake rather than Google Cloud, and I would bring that production cloud data platform experience to Halden Pay's team.

I would welcome the opportunity to discuss how my experience building dependable, reusable data platform capabilities could contribute to Halden Pay. Thank you for your consideration.

Kind regards,
Lena Varga-Holm