Lena Varga-Holm

Senior Data Engineer | Data Platform

Email: applicant.inbox@example.com | Phone: +44 7700 900123

PROFESSIONAL SUMMARY

Senior Data Engineer with experience owning data platform ingestion and orchestration, building production data pipelines, and operating cloud data warehouses. Delivered streaming and change data capture pipelines, migrated warehouse transformations to dbt, and improved platform reliability, build times, and operating costs. Hands-on with Python, SQL, Airflow, Kafka, Snowflake, dbt, and AWS.

TECHNICAL SKILLS

Languages: Python, SQL
Data platform and orchestration: Airflow, Kafka, Kafka Connect, Debezium, dbt
Warehousing and cloud: Snowflake, AWS MSK, AWS EMR, Amazon S3, Redshift, Glue Data Catalog, Snowpipe Streaming
Quality and delivery: Great Expectations, pytest, GitHub Actions, Terraform

PROFESSIONAL EXPERIENCE

Senior Data Engineer | Lukovac Freight Hub d.o.o. | 2022–Present

- Own ingestion and orchestration for the data platform using Airflow, Snowflake, dbt, and Kafka.
- Participate in the data platform on-call rotation, one week in six.
- Mentor two mid-level engineers and run data platform office hours.
- Reduced Snowflake spend from approximately 38k to 24k per month in 2024 through warehouse sizing, auto-suspend, and removal of unused tables.

Data Engineer | Kovalent Pay | 2019–2022

- Developed PySpark data processing on AWS EMR and designed Redshift warehouse solutions.
- Applied Great Expectations for data quality.

SELECTED PROJECTS

Shipment Event Streaming Pipeline | Lukovac Freight Hub | 2023

- Designed and built a shipment tracking pipeline ingesting courier-partner webhooks into Kafka and Snowflake through Kafka Connect and Snowpipe Streaming.
- Processed approximately 30 million events per day with end-to-end latency under two minutes at p95.
- Added a dead-letter topic and replay tooling to reprocess failed events without data loss.
- Technologies: Kafka, Kafka Connect, Snowflake, Python, Terraform, AWS MSK.

Change Data Capture from Postgres | Lukovac Freight Hub | 2024

- Replaced hourly full-table extracts from the orders database with Debezium CDC into Kafka.
- Reduced replica CPU during loads from approximately 70% to 25%.
- Technologies: Debezium, Kafka, PostgreSQL, Airflow, Snowflake.

dbt Migration of the Warehouse | Lukovac Freight Hub | 2022–2023

- Migrated approximately 300 stored procedures into dbt models with tests and documentation.
- Introduced incremental models and slim CI using state:modified in GitHub Actions.
- Reduced the nightly warehouse build from 5 hours 10 minutes to 1 hour 35 minutes and made lineage visible to analysts.
- Technologies: dbt, Snowflake, SQL, GitHub Actions.

Transaction Reconciliation Engine | Kovalent Pay | 2020–2021

- Built PySpark batch jobs on EMR to match bank settlement files against internal ledgers.
- Reduced manual finance-team reconciliation from approximately three days to four hours per

month-end.
- Technologies: PySpark, AWS EMR, S3, Glue Data Catalog, Redshift, Great Expectations.

Open-source Contribution | Airflow provider package | 2023

- Fixed a pagination bug and added tests for a community Airflow provider for a SaaS CRM API; two pull requests were merged.
- Technologies: Python, Airflow, pytest.

EDUCATION

MSc in Computer Science | Northgate Technical University
Thesis: Incremental materialized views

CERTIFICATION

AWS Certified Data Engineer - Associate | May 2024

COMMUNITY

- Teach SQL for beginners at a local women-in-tech meetup, approximately four evenings per year since 2021.
- Presented "Data contracts without the hype" at a regional data meetup in 2024.