Dear Hiring Manager,

I am applying for the Senior Data Engineer – Data Platform role at Halden Pay. In my current role at Lukovac Freight Hub, I own ingestion and orchestration for a production data platform using Airflow, Snowflake, dbt, and Kafka. Building shared platform capabilities and reliable pipelines is central to my work, which aligns closely with Halden Pay's focus on reusable frameworks for analysts, data scientists, and product teams.

At Lukovac, I designed a shipment event streaming pipeline processing approximately 30 million events per day, with end-to-end latency under two minutes at p95, and added dead-letter and replay tooling to support recovery without data loss. I also migrated approximately 300 stored procedures into tested and documented dbt models, reducing the nightly warehouse build from 5h10 to 1h35. In 2024, I reduced monthly Snowflake spend from approximately USD 38,000 to USD 24,000 through warehouse sizing, auto-suspend, and removing unused tables. These projects combine platform operation, data quality, and practical improvements to reliability and efficiency.

Earlier, at Kovalent Pay, I built PySpark jobs on AWS EMR for a transaction reconciliation engine that reduced finance's manual month-end reconciliation work from approximately three days to four hours. My experience with AWS-managed data services and Snowflake provides a strong foundation in comparable cloud data environments, and I hold the AWS Certified Data Engineer – Associate certification. I also mentor colleagues, teach SQL at a local women-in-tech meetup, and have contributed fixes and tests to an open-source Airflow provider.

Halden Pay's emphasis on ownership, reducing unnecessary complexity, and building a shared data platform is a strong match for how I approach data engineering. I would welcome the opportunity to discuss how my experience could contribute to the Data Platform team.

Kind regards,
Lena Varga-Holm