Dear Hiring Team at Halden Pay,

I am applying for the Senior Data Engineer – Data Platform position. Over nine years, I have built and operated production data platforms and cloud warehouses across logistics, e-commerce and payments. Halden Pay's focus on a shared platform for product, analytics and data science teams is a strong match for the work I do: building reliable, reusable data systems that help teams make good use of trusted data.

At Lukovac Freight Hub, I own ingestion and orchestration across a platform with around 140 Airflow DAGs, and I migrated warehouse transformations to more than 420 tested dbt models. That work reduced the nightly warehouse build from 5 hours 10 minutes to 1 hour 50 minutes. I have also introduced data contracts with JSON Schema and CI checks, and built a Kafka-based shipment-event pipeline processing around 30 million events per day. I also reduced Snowflake spend from approximately 38k to 24k per month through warehouse sizing, auto-suspend and removal of unused tables. Together, these projects demonstrate my experience designing, building and operating a production data platform, and implementing data quality controls and reusable platform foundations.

My experience in payments is directly relevant to Halden Pay. At Kovalent Pay, I built PySpark jobs on AWS EMR to reconcile bank settlement files against internal ledgers, processing about 1.2 TB of raw files daily. The work reduced month-end manual reconciliation from around three days to four hours. I also designed a Redshift star-schema warehouse, applying dimensional modelling across more than 30 fact and dimension tables for finance and risk teams. I also wrote an internal Python library for S3 and the Glue Data Catalog that was adopted by four teams, and established Great Expectations checks for critical finance tables. Alongside hands-on AWS experience with S3, EMR, Glue, Lambda, IAM and Docker-based engineering workflows, I hold the AWS Certified Data Engineer – Associate certification.

I value clear data contracts, well-tested pipelines and open-source collaboration. I am particularly interested in building reusable frameworks rather than one-off pipelines. I have applied this approach through shared data contracts, tested dbt models, CDC and replay tooling, and internal Python libraries used by multiple teams. I have contributed two merged pull requests to an Airflow provider, and I would welcome the opportunity to bring this practical, ownership-focused approach to Halden Pay's shared Data Platform, building reusable and well-tested foundations in collaboration with product teams and data users. I would be glad to discuss how my experience could contribute.

Kind regards,

Lena Varga-Holm
lenavh.dev