Dear Halden Pay Hiring Team,

I am applying for the Senior Data Engineer – Data Platform role. Across nine years in data engineering, including work in payments, I have built and operated production data platforms and warehouses that support finance, risk, analytics and product teams. I am particularly drawn to Halden Pay's focus on reusable platform capabilities and its values of ownership, collaboration and moving quickly by reducing unnecessary complexity.

At Lukovac Freight Hub, I own the ingestion layer of a production data platform with around 140 Airflow DAGs and have helped make its warehouse more reliable and efficient. I moved transformation logic into dbt, with tests across 420-plus models, reducing the nightly build from 5 hours 10 minutes to 1 hour 35 minutes. I also introduced data contracts with CI checks, reducing broken-pipeline incidents, and built a Kafka-based streaming pipeline processing around 30 million events per day. These projects have involved designing, building, and operating production data-platform capabilities, improving data quality and making platform interfaces more dependable for the teams that use them.

My experience at Kovalent Pay also gives me a direct connection to Halden Pay's payments work. I built PySpark jobs on AWS EMR to reconcile bank settlement files with internal ledgers, processing about 1.2 TB per day and reducing month-end manual reconciliation from around three days to four hours. I designed a Redshift dimensional warehouse using a star schema with more than 30 fact and dimension tables, used by finance and risk, and added data quality checks to critical finance tables. This work required writing production Python and SQL for large datasets and optimising data processing and warehouse performance; at Lukovac, I also reduced the Snowflake bill from approximately $38K to $24K per month through warehouse sizing, auto-suspend, and removal of unused tables. My cloud experience is with AWS, including managed data services such as EMR, S3 and Redshift; I would bring that production experience alongside my Python, SQL, Airflow and dbt background.

I value open-source practices and knowledge-sharing: I have contributed fixes and tests to a community Airflow provider, mentor two data engineers, and have spoken at a regional data meetup about data contracts. I would welcome the opportunity to discuss how my experience building reliable, reusable data platforms could support Halden Pay's Data Platform team.

Kind regards,
Lena Varga-Holm
lenavh.dev