LENA VARGA-HOLM

Senior Data Engineer

Rijeka, Croatia (open to remote, EU time zones) | lena.vargaholm@example.com | +385 55 555 0101
github.com/lvh-data-example | lenavh.dev

PROFILE

Data engineer with 9 years of experience building batch and streaming data platforms for e-commerce and payments. I specialise in reliable, well-tested production data platforms, clear data contracts, dimensional warehouses, and batch and streaming pipelines used by analytics and business teams.

EXPERIENCE

Senior Data Engineer - Lukovac Freight Hub d.o.o. (logistics marketplace), Rijeka
03/2022 - present

- Own ingestion and orchestration for the company data platform using Airflow, Snowflake, dbt and Kafka.
- Migrated approximately 300 stored procedures into dbt models with tests and documentation; reduced the nightly warehouse build from 5h10 to 1h50.
- Introduced data contracts between product teams and the data team, using CI checks to improve pipeline reliability.
- Built a streaming pipeline for shipment tracking events using Kafka, Kafka Connect and Snowpipe Streaming (~30M events/day), with under 2 minutes p95 end-to-end latency and dead-letter and replay tooling.
- Replaced hourly full-table extracts from PostgreSQL with Debezium CDC into Kafka, reducing replica CPU during loads from ~70% to ~25%.
- Mentor two mid-level data engineers; run the weekly data platform office hours.
- Participate in the data-platform on-call rotation, one week in six.
- Optimised Snowflake warehouse usage through warehouse sizing, auto-suspend and removal of unused tables, reducing spend from approximately $38K/month to $24K/month.

Data Engineer - Kovalent Pay (payments startup), Zagreb *09/2019 - 02/2022*

- Built PySpark jobs on AWS EMR for transaction reconciliation (daily ~1.2 TB of raw files from S3).
- Designed the Redshift warehouse using dimensional modelling and a star schema, with 30+ fact and dimension tables used by finance and risk.
- Set up data quality checks with Great Expectations on critical finance tables.
- Wrote internal Python library for S3/Glue catalog access used by 4 teams.

Junior Data Engineer / BI Developer - Adria Retail Group, Split *06/2017 - 08/2019*

- Maintained SSIS packages and T-SQL ETL into an on-prem SQL Server warehouse.
- Built Power BI reports for store operations (40+ stores).
- Migrated 60 nightly SSIS jobs to Python scripts scheduled with cron.

SELECTED ENGINEERING CONTRIBUTIONS

- Open-source contribution: fixed a pagination bug and added tests to an Airflow provider package; 2 merged pull requests (2023).
- Presented 'Data contracts without the hype' at a regional data meetup (2024).

- Teach SQL to beginners at a local women-in-tech meetup, approximately four evenings per year since 2021.

EDUCATION

MSc Computer Science (Data Science track) - Northgate Technical University, 2017
Thesis: Incremental materialized views for analytical queries

SKILLS

Languages: Python, SQL, Bash
Data Platforms & Modelling: Airflow, dbt, Spark (PySpark), Kafka, Kafka Connect, Snowflake, Redshift, PostgreSQL, Great Expectations, dimensional modelling, data contracts, data quality, CDC/Debezium
Cloud & infra: AWS (S3, EMR, Glue, Lambda, IAM), Docker, Terraform (basic), GitHub Actions
BI: Power BI, Looker (basic)
Certification: AWS Certified Data Engineer – Associate (May 2024)

LANGUAGES

Croatian (native), English (C1), German (B1)