Dear Halden Pay Hiring Team,

I am applying for the Senior Data Engineer – Data Platform role. Over nine years in data engineering, I have designed, built and operated production data platforms across payments, e-commerce and logistics. In my current role, I own the ingestion layer of a platform with around 140 Airflow DAGs and have helped make its pipelines faster, more reliable and easier for teams to use.

At Lukovac Freight Hub, I migrated warehouse transformations to dbt, adding tests across more than 420 models and reducing the nightly run from 5 hours 10 minutes to 1 hour 35 minutes. I also introduced data contracts with CI checks, reducing broken-pipeline incidents from around nine per month to two. Alongside batch workloads, I built a Kafka-based streaming pipeline processing approximately 30 million shipment events per day. These experiences align closely with owning a shared platform and building reusable, dependable capabilities for analytics and product teams.

My earlier work at Kovalent Pay gave me direct experience in payments data. I built PySpark jobs on AWS EMR to reconcile around 1.2 TB of daily files against internal ledgers, reducing month-end reconciliation from approximately three days to four hours. I also designed a Redshift dimensional warehouse using a star schema, with more than 30 fact and dimension tables, for finance and risk and introduced quality checks on critical finance tables. My cloud experience is with AWS and Snowflake, alongside Python, SQL, Airflow, dbt, Docker and CI/CD. I have also optimised Snowflake warehouse performance and cost through warehouse sizing, auto-suspend settings and unused-table retention, reducing monthly spend from approximately $38K to $24K in 2024.

I value clear data contracts, well-tested pipelines, data quality standards and open-source practices. I have contributed fixes and tests to a community Airflow provider, and I mentor colleagues and run data platform office hours. Halden Pay’s focus on shared platform capabilities, ownership and reducing unnecessary complexity is a strong match for how I approach data engineering. I would welcome the opportunity to discuss how I could contribute to the Data Platform team.

Kind regards,

Lena Varga-Holm