Dear Halden Pay Hiring Team,

I am applying for the Senior Data Engineer – Data Platform role. Over nine years in data engineering, I have built and operated production data platforms for logistics, e-commerce and payments, with a focus on dependable pipelines, reusable capabilities and data that business teams can trust.

At Lukovac Freight Hub, I own the ingestion and orchestration layer, operating around 140 Airflow DAGs and a Snowflake warehouse with more than 420 dbt models. I migrated approximately 300 stored procedures into dbt, added tests and CI, and reduced the nightly build from 5 hours 10 minutes to 1 hour 50 minutes. I also introduced JSON Schema data contracts and CI checks that helped reduce broken-pipeline incidents from about nine to two per month. This work aligns closely with your focus on a shared platform and reusable frameworks.

I have also built streaming and batch systems at scale: a Kafka shipment-event pipeline processing around 30 million events per day, and PySpark reconciliation jobs on AWS EMR processing approximately 1.2 TB of daily files. The reconciliation work cut manual month-end effort from about three days to four hours. My experience includes AWS, Snowflake, Redshift, Python, SQL, Docker and GitHub Actions, as well as dimensional modelling, data quality testing and mentoring engineers.

I am interested in Halden Pay's Data Platform role because it brings platform engineering together with the needs of analytics, data science and product teams. I would welcome the opportunity to discuss how my experience building reliable, reusable data systems could contribute to the team.

Kind regards,
Lena Varga-Holm
lenavh.dev